

train_diffusion

January 28, 2026

1 Embryo Diffusion — Next-Frame Prediction (silver F0)

This notebook trains a **context-conditioned diffusion model** to predict the next frame from the last k context frames. Training optimizes **noise-prediction MSE** (with optional LPIPS and inversion regularizer), supports **cosine LR decay with warmup**, and includes **newest-frame corruption** for robustness.

What it does - Builds a cached index from *silver* frames and phase CSVs. - Splits by **embryo ID** to avoid leakage. - Trains a UNet to predict noise on the target frame, conditioned on the k previous frames and an optional absolute frame index embedding. - Uses DDPM for training schedule and DDIM for fast sampling/inversion. - Optionally uses external **phase classifier guidance** during sampling only.

Project structure (expected)

```
thesis/  
  data/  
    embryo_dataset_silver/  
      <plane_id>/<embryo_id>/  
        00001.jpeg, 00002.jpeg, ...  
        <embryo_id>_phases.csv  
  classifier/  
    embryo_phase_models/  
  diffusion/  
    train_diffusion.ipynb  
    embryo_diffusion_models/  
    diffusion_samples/
```

Notes - Presence filtering is **not** run here; *silver* is already filtered. - Training pairs are adjacent in time: context $[t - k + 1..t] \rightarrow$ target $t + 1$. - External classifier guidance (if enabled) is used **only at sampling**, not in the training loss.

1.1 Notebook map

This notebook is long, so here is the high-level flow: 1. **Configuration & paths** — dataset locations, training hyperparameters, and guidance toggles. 2. **Imports + reproducibility** — libraries, seeds, and the phase label set. 3. **Indexing** — build a frame index from the *silver* dataset (cached on disk). 4. **Splits** — split **by embryo** to prevent leakage across train/val/test. 5. **Model & diffusion schedule** — UNet backbone and DDPM/DDIM schedulers. 6. **Dataset & loaders** — preprocessing utilities + context $\rightarrow$ target dataset. 7. **Sampling helpers** — model

init + DDIM sampling/guidance utilities. 8. **Training** — loss, metrics, checkpoints, and sample visualizations. 9. **Inference utilities** — ad hoc sampling and multi-step rollouts.

1.2 1. Configuration

This cell defines all paths, data splits, training hyperparameters, and guidance switches. If you opened the notebook from the repo root, it automatically points to the diffusion subfolder so paths stay correct.

```
[1]: # ===== CONFIG (edit as needed) =====
# This cell centralizes *all* the knobs for data paths, training, and sampling.
from pathlib import Path
import os

# DEBUG MODE: Set to True for quick testing with minimal data
DEBUG_MODE = False
VERBOSE = True # extra logging for dataset/split diagnostics

# Project root resolution: if you open from repo root, auto-jump into /diffusion
PROJECT_ROOT = Path.cwd().resolve()
if PROJECT_ROOT.name != 'diffusion':
    # try to infer if opened from repo root
    if (PROJECT_ROOT / 'diffusion').exists():
        PROJECT_ROOT = PROJECT_ROOT / 'diffusion'

# Data roots
DATA_ROOT = PROJECT_ROOT.parent / 'data'
SILVER_ROOT = DATA_ROOT / 'embryo_dataset_silver'
FOCAL_PLANES = ['F0'], 'F-15', 'F-30', 'F-45', 'F15', 'F30', 'F45']

# Context window (k previous frames) for next-frame prediction
CONTEXT_K = 7 # Sharma-style fixed window (F=7)

# Time conditioning with absolute frame index embedding
USE_TIME_MAP = True
MAX_FRAME_INDEX = 500 # safe upper bound (dataset max ~450)
TIME_EMB_DIM = 16 # small, cheap

# Outputs
OUT_MODELS_DIR = PROJECT_ROOT / 'embryo_diffusion_models'
OUT_SAMPLES_DIR = PROJECT_ROOT / 'diffusion_samples'
CACHE_DIR = PROJECT_ROOT / 'cache'
for d in (OUT_MODELS_DIR, OUT_SAMPLES_DIR, CACHE_DIR):
    d.mkdir(parents=True, exist_ok=True)

# Training hyperparameters
IMG_SIZE = 128 # images are 500x500 in dataset; we center-crop &
↳resize
```

```

BATCH_SIZE = 16
LR = 2e-4 # Reduced from 2e-4 to prevent gradient explosion
USE_LR_SCHEDULER = True # LROnPlateau only
LR_MIN = 1e-7 # higher floor to avoid over-decay
EPOCHS = 1000 if not DEBUG_MODE else 4 # Reduced epochs for debug
GRAD_CLIP = 1.0 # Gradient clipping to prevent NaN
NUM_WORKERS = 8
VAL_SPLIT = 0.10
TEST_SPLIT = 0.10
SPLIT_SEED = 42

# Diffusion process
TIMESTEPS = 1000 # training timesteps
BETA_START, BETA_END = 1e-4, 2e-2

# DDIM sampling/inversion
USE_DDIM_SAMPLING = True
DDIM_STEPS = 150
DDIM_ETA = 0.0

# Inversion training (optional regularizer)
INVERSION_TRAINING = True
INVERSION_LAMBDA = 0.1
INVERSION_STEPS = 10
INVERSION_EVERY_N = 16
INVERSION_MAX_BATCH = 8
USE_GRAD_CHECKPOINTING = False # disable to avoid CheckpointError in PyTorch 2.
↪x

# Sampling visualization
SAMPLE_PLANES = ['F0', 'F-15', 'F30']
SAMPLES_PER_PLANE = 3

# Data augmentation (applied identically to context stack + target)
AUGMENT = True
AUGMENT_HFLIP = 0.5
AUGMENT_VFLIP = 0.1
AUGMENT_MAX_ROT = 5.0
AUGMENT_BRIGHTNESS = 0.1
AUGMENT_CONTRAST = 0.1

# Loss weights
LPIPS_LOSS_WEIGHT = 0.05
XO_L1_WEIGHT = 0.1 # image-space L1 on x0_pred
XO_SSIM_WEIGHT = 0.1 # SSIM-style term (1 - SSIM)

# Validation metrics (SSIM/LPIPS)

```

```

COMPUTE_VAL_METRICS = True
VAL_METRIC_SAMPLES = 256 if not DEBUG_MODE else 4 # number of validation pairs
↳evaluated per epoch
VAL_METRIC_MAX_BATCHES = 4 # safety cap for validation metric computation

# Debug overrides
if DEBUG_MODE:
    NUM_WORKERS = 2
    BATCH_SIZE = 4
    DDIM_STEPS = 20
    INVERSION_STEPS = 8
    INVERSION_EVERY_N = 10
    INVERSION_MAX_BATCH = 2
    SAMPLE_PLANES = ['F0']
    SAMPLES_PER_PLANE = 1

print('Project root:', PROJECT_ROOT)
print('Silver root:', SILVER_ROOT)
print('Focal planes:', FOCAL_PLANES)
print(f'Context window k={CONTEXT_K}')
print(f'Splits: val={VAL_SPLIT:.2f}, test={TEST_SPLIT:.2f}, seed={SPLIT_SEED}')
print(f'DDIM: steps={DDIM_STEPS}, eta={DDIM_ETA}, enabled={USE_DDIM_SAMPLING}')
print(f'Inversion training: {INVERSION_TRAINING} (lambda={INVERSION_LAMBDA},
↳steps={INVERSION_STEPS}, every={INVERSION_EVERY_N},
↳max_batch={INVERSION_MAX_BATCH})')
print(f'Grad checkpointing: {USE_GRAD_CHECKPOINTING}')
if DEBUG_MODE:
    print('  DEBUG_MODE is ON - using reduced data for testing')
for plane in FOCAL_PLANES:
    plane_dir = SILVER_ROOT / plane
    assert plane_dir.exists(), f"Silver {plane} folder not found: {plane_dir}"

```

```

Project root: /mnt/c/Users/ioann/OneDrive/Documents/Projects/thesis/diffusion
Silver root:
/mnt/c/Users/ioann/OneDrive/Documents/Projects/thesis/data/embryo_dataset_silver
Focal planes: ['F0']
Context window k=7
Splits: val=0.10, test=0.10, seed=42
DDIM: steps=150, eta=0.0, enabled=True
Inversion training: True (lambda=0.1, steps=10, every=16, max_batch=8)
Grad checkpointing: False

```

1.3 2. Imports and small utilities

The notebook relies on **PyTorch** for training, **diffusers** for the UNet + schedulers, and **LPIPS** for perceptual validation loss.

```
[2]: # Imports & utils
import math, random, glob, re, os
import numpy as np
import pandas as pd
from PIL import Image
from tqdm import tqdm
import matplotlib.pyplot as plt

# PyTorch core
import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import Dataset, DataLoader
from torchvision.transforms import functional as TF

# Perceptual metric
from lpips import LPIPS

# Diffusers library for UNet and schedulers
from diffusers import UNet2DModel, DDPM Scheduler, DDIM Scheduler

# TQDM helper: force ETA/remaining to show consistently
TQDM_BAR_FORMAT = "{l_bar}{bar} {n_fmt}/{total_fmt} [{elapsed}<{remaining},
↳{rate_fmt}]"

def tqdm_bar(iterable, **kwargs):
    """Convenience wrapper with consistent progress-bar formatting."""
    return tqdm(iterable, leave=True, bar_format=TQDM_BAR_FORMAT, **kwargs)
```

1.4 3. Reproducibility and label setup

We seed all relevant RNGs for reproducibility, pick the compute device, and define the ordered phase labels used throughout the notebook.

```
[3]: # Initialize seed and device (must be before using PHASE_LABELS in later cells)
def seed_everything(seed=42):
    """Make results reproducible across Python/NumPy/PyTorch."""
    random.seed(seed)
    np.random.seed(seed)
    torch.manual_seed(seed)
    torch.cuda.manual_seed_all(seed)
    torch.backends.cudnn.deterministic = True
    torch.backends.cudnn.benchmark = False

seed_everything(42)
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print('Using device:', device)
```

```

# Ordered list of morphokinetic phase labels (target classes)
PHASE_LABELS = [
    ↪['pPB2', 'pPNa', 'pPNf', 'p2', 'p3', 'p4', 'p5', 'p6', 'p7', 'p8', 'p9+', 'pM', 'pSB', 'pB', 'pEB', 'pHB']
NUM_PHASES = len(PHASE_LABELS)

# Plane ID mapping (only uses planes in FOCAL_PLANES)
PLANE_TO_ID = {plane: i for i, plane in enumerate(FOCAL_PLANES)}
ID_TO_PLANE = {i: plane for plane, i in PLANE_TO_ID.items()}
NUM_PLANES = len(FOCAL_PLANES)

```

Using device: cuda

1.5 4. Build a frame index (cached)

We scan the *silver* dataset and build a table of frame paths + phase labels. This is cached to CSV so you do it once.

```

[4]: # Build frame index from all focal planes: phase CSV lives next to frames
# Output columns: embryo_id, plane, frame_num (1-based), path, phase_id
INDEX_CACHE = CACHE_DIR / 'silver_all_planes_frame_index.csv'

def build_index_from_silver():
    rows = []
    # For each focal plane
    for plane in FOCAL_PLANES:
        plane_root = SILVER_ROOT / plane
        embryo_ids = sorted([d.name for d in plane_root.iterdir() if d.
↪is_dir()])
        # For each embryo id included in focal plane
        for eid in tqdm(embryo_ids, desc=f'Indexing {plane}'):
            folder = plane_root / eid
            ann_path = folder / f"{eid}_phases.csv"
            if not ann_path.exists():
                continue
            # Use .jpeg files (silver dataset uses .jpeg)
            img_files = sorted(glob.glob(str(folder / '*.jpeg')))
            if not img_files:
                continue
            # Extract frame numbers from filenames (00037.jpeg → 37)
            frame_numbers = {}
            # For each frame image file
            for f in img_files:
                basename = os.path.basename(f)
                m = re.search(r'(\d+)\.jpeg$', basename)
                if m:
                    frame_num = int(m.group(1))
                    frame_numbers[frame_num] = f

```

```

    # read CSV (phase, start, end) with 1-based frame indices
    df = pd.read_csv(ann_path, header=None,
names=['phase', 'start', 'end'], sep='[;\s]+', engine='python')
    # Phase must start with p, e.g., pPNA not tPNA
    def phase_to_id(ph: str):
        ph = ph.strip()
        ph = ph.replace('t', 'p', 1) if ph.startswith('t') else ph
    accept t* or p*
        return PHASE_LABELS.index(ph)
    label_spans = []
    for _, r in df.iterrows():
        try:
            pid = phase_to_id(r['phase'])
            label_spans.append((int(r['start']), int(r['end']), pid))
        except Exception:
            pass
    # assign phase per frame using actual frame numbers from filenames
    for frame_num, path in frame_numbers.items():
        pid = None
        for s, e, p in label_spans:
            if s <= frame_num <= e: # frame_num is 1-based, matches CSV
                pid = p
                break
        if pid is None:
            continue
        rows.append({'embryo_id': eid, 'plane': plane, 'frame_num':
frame_num, 'path': path, 'phase_id': pid})
    return pd.DataFrame(rows)

```

1.5.1 4.1 Index caching

We cache the full frame index to speed up restarts. Delete the cache file if you change the dataset layout.

```

[5]: if INDEX_CACHE.exists():
        index_df = pd.read_csv(INDEX_CACHE)
    else:
        index_df = build_index_from_silver()
        index_df.to_csv(INDEX_CACHE, index=False)

    print('Indexed frames (full):', len(index_df))
    print('Planes:', sorted(index_df['plane'].unique().tolist()))
    index_df.head()

```

Indexing F0: 100% | 704/704 [00:06<00:00, 108.79it/s]

Indexed frames (full): 289675

Planes: ['F0']

```
[5]: embryo_id plane frame_num \
0    AA83-7    F0          5
1    AA83-7    F0          6
2    AA83-7    F0          7
3    AA83-7    F0          8
4    AA83-7    F0          9

path phase_id
0 /mnt/c/Users/ioann/OneDrive/Documents/Projects... 0
1 /mnt/c/Users/ioann/OneDrive/Documents/Projects... 0
2 /mnt/c/Users/ioann/OneDrive/Documents/Projects... 0
3 /mnt/c/Users/ioann/OneDrive/Documents/Projects... 0
4 /mnt/c/Users/ioann/OneDrive/Documents/Projects... 0
```

1.6 5. Train/val/test splits (by embryo)

We split **by embryo ID**, not by frame, to avoid leaking information across sets. Splits are cached on disk.

```
[6]: # Build embryo splits (no embryo appears in more than one split)
SPLIT_CACHE_STEM = (
    f"splits_embryo_phase_ctxk{CONTEXT_K}_"
    f"seed{SPLIT_SEED}_v{VAL_SPLIT}_t{TEST_SPLIT}"
    + ("_debug" if DEBUG_MODE else "")
)
TRAIN_SPLIT_CACHE = CACHE_DIR / f"{SPLIT_CACHE_STEM}_train.csv"
VAL_SPLIT_CACHE = CACHE_DIR / f"{SPLIT_CACHE_STEM}_val.csv"
TEST_SPLIT_CACHE = CACHE_DIR / f"{SPLIT_CACHE_STEM}_test.csv"

def stratified_split(ids: list[str], labels: dict[str, int], val_frac: float,
    test_frac: float, seed: int):
    """Stratified split by dominant phase label per embryo."""
    rng = np.random.RandomState(seed)
    train_ids, val_ids, test_ids = [], [], []
    by_label: dict[int, list[str]] = {}
    for eid in ids:
        by_label.setdefault(labels[eid], []).append(eid)
    for label, group in by_label.items():
        group = list(group)
        rng.shuffle(group)
        n = len(group)
        n_test = int(round(n * test_frac))
        n_val = int(round(n * val_frac))
        n_test = min(n_test, n)
        n_val = min(n_val, n - n_test)
        test_ids.extend(group[:n_test])
        val_ids.extend(group[n_test:n_test + n_val])
```



```

        train_ids.extend(group[n_test + n_val:])
    return set(train_ids), set(val_ids), set(test_ids)

# In DEBUG mode, sample only a few embryos
index_for_pairs = index_df.copy()
if DEBUG_MODE:
    unique_embryos = index_for_pairs["embryo_id"].unique()
    if len(unique_embryos) > 3:
        debug_embryos = np.random.choice(unique_embryos, 3, replace=False)
        index_for_pairs = index_for_pairs[index_for_pairs["embryo_id"].
↪isin(debug_embryos)]
        print(f"DEBUG_MODE: Using only {len(debug_embryos)} embryos for_
↪testing")
    else:
        debug_embryos = unique_embryos

print("Indexed frames available:", len(index_for_pairs))
print("Planes in index:", sorted(index_for_pairs["plane"].unique().tolist()))

def _debug_split(ids: list[str]):
    """Small deterministic split for quick debug runs."""
    ids = list(ids)
    rng = np.random.RandomState(SPLIT_SEED)
    rng.shuffle(ids)
    if len(ids) >= 3:
        return set(ids[2:]), {ids[1]}, {ids[0]}
    if len(ids) == 2:
        return {ids[1]}, {ids[0]}, set()
    if len(ids) == 1:
        return {ids[0]}, set(), set()
    return set(), set(), set()

if TRAIN_SPLIT_CACHE.exists() and VAL_SPLIT_CACHE.exists() and TEST_SPLIT_CACHE.
↪exists():
    train_df = pd.read_csv(TRAIN_SPLIT_CACHE)
    val_df = pd.read_csv(VAL_SPLIT_CACHE)
    test_df = pd.read_csv(TEST_SPLIT_CACHE)
else:
    # Dominant phase per embryo for stratification
    embryo_phase = (
        index_for_pairs.groupby("embryo_id")["phase_id"]
        .agg(lambda s: s.value_counts().idxmax())
        .to_dict()
    )
    embryo_ids = list(embryo_phase.keys())
    if DEBUG_MODE:
        train_ids, val_ids, test_ids = _debug_split(embryo_ids)

```

```

else:
    train_ids, val_ids, test_ids = stratified_split(embryo_ids,
↪embryo_phase, VAL_SPLIT, TEST_SPLIT, SPLIT_SEED)
    train_df = pd.DataFrame({"embryo_id": list(train_ids)})
    val_df = pd.DataFrame({"embryo_id": list(val_ids)})
    test_df = pd.DataFrame({"embryo_id": list(test_ids)})
    train_df.to_csv(TRAIN_SPLIT_CACHE, index=False)
    val_df.to_csv(VAL_SPLIT_CACHE, index=False)
    test_df.to_csv(TEST_SPLIT_CACHE, index=False)

train_ids = set(train_df["embryo_id"].tolist())
val_ids = set(val_df["embryo_id"].tolist())
test_ids = set(test_df["embryo_id"].tolist())

# DEBUG safeguard: ensure we have a non-empty val set for metrics/samples
if DEBUG_MODE:
    all_ids = sorted(set(index_for_pairs["embryo_id"].unique().tolist()))
    if len(val_ids) == 0 and len(all_ids) >= 2:
        train_ids, val_ids, test_ids = _debug_split(all_ids)
        train_df = pd.DataFrame({"embryo_id": list(train_ids)})
        val_df = pd.DataFrame({"embryo_id": list(val_ids)})
        test_df = pd.DataFrame({"embryo_id": list(test_ids)})
        train_df.to_csv(TRAIN_SPLIT_CACHE, index=False)
        val_df.to_csv(VAL_SPLIT_CACHE, index=False)
        test_df.to_csv(TEST_SPLIT_CACHE, index=False)

print(f"Splits cached: {TRAIN_SPLIT_CACHE.name}, {VAL_SPLIT_CACHE.name},
↪{TEST_SPLIT_CACHE.name}")
print(f"Train embryos: {len(train_ids)} | Val embryos: {len(val_ids)} | Test
↪embryos: {len(test_ids)}")

if VERBOSE:
    def _split_stats(name: str, ids: set[str]):
        df_split = index_for_pairs[index_for_pairs["embryo_id"].isin(ids)]
        frames = len(df_split)
        sequences = df_split.groupby(["embryo_id", "plane"]).ngroups
        planes = df_split["plane"].nunique()
        embryos = df_split["embryo_id"].nunique()
        print(f"[{name}] embryos={embryos} | planes={planes} |
↪sequences={sequences} | frames={frames}")
    _split_stats("train", train_ids)
    _split_stats("val", val_ids)
    _split_stats("test", test_ids)

```

Indexed frames available: 289675

Planes in index: ['F0']

Splits cached: splits_embryo_phase_ctxk7_seed42_v0.1_t0.1_train.csv,

```
splits_embryo_phase_ctxk7_seed42_v0.1_t0.1_val.csv,
splits_embryo_phase_ctxk7_seed42_v0.1_t0.1_test.csv
Train embryos: 565 | Val embryos: 69 | Test embryos: 69
[train] embryos=565 | planes=1 | sequences=565 | frames=230158
[val] embryos=69 | planes=1 | sequences=69 | frames=30890
[test] embryos=69 | planes=1 | sequences=69 | frames=28627
```

1.7 6. Diffusion schedules

We use **DDPM** for training noise schedules and **DDIM** for fast sampling and inversion.

```
[7]: # Diffusion schedules using diffusers

# Training schedule (DDPM)
scheduler = DDPMScheduler(
    num_train_timesteps=TIMESTEPS,
    beta_start=BETA_START,
    beta_end=BETA_END,
    beta_schedule='scaled_linear',
    variance_type='fixed_small',
)

# Sampling schedule (DDIM for fast inference)
ddim_scheduler = DDIMScheduler(
    num_train_timesteps=TIMESTEPS,
    beta_start=BETA_START,
    beta_end=BETA_END,
    beta_schedule='scaled_linear',
    clip_sample=False,
    set_alpha_to_one=False,
)

# Pre-compute noise schedule for training
betas = scheduler.betas.to(device)
alphas = (1 - betas).to(device)
alphas_cumprod = torch.cumprod(alphas, dim=0).to(device)
sqrt_alphas_cumprod = torch.sqrt(alphas_cumprod)
sqrt_one_minus_alphas_cumprod = torch.sqrt(1.0 - alphas_cumprod)

# Compute posterior variance for DDPM sampling
# posterior_var_t = (1 - alpha_cumprod_{t-1}) / (1 - alpha_cumprod_t) * beta_t
alphas_cumprod_prev = torch.cat([torch.ones(1, device=device), alphas_cumprod[:
    ↪-1]])
posterior_var = (1 - alphas_cumprod_prev) / (1 - alphas_cumprod) * betas
posterior_var = torch.clamp(posterior_var, min=1e-20) # avoid log(0)
```

1.8 7. Model: context-conditioned UNet

We predict the *noise* added to the target frame while conditioning on a stack of k previous frames and (optionally) an absolute frame index embedding.

```
[8]: # Simplified diffusion model using diffusers UNet2DModel (context-conditioned)

class DiffusionCond(nn.Module):
    """Diffusion model conditioned on a context stack of k previous frames."""
    def __init__(self, context_k: int = CONTEXT_K, use_time_map: bool = USE_TIME_MAP):
        super().__init__()
        self.context_k = context_k
        self.use_time_map = use_time_map

        # Learned absolute frame index embedding (Sharma-style elapsed time
        # proxy)
        if self.use_time_map:
            self.frame_emb = nn.Embedding(MAX_FRAME_INDEX + 1, TIME_EMB_DIM)
            self.frame_proj = nn.Linear(TIME_EMB_DIM, 1)
        else:
            self.frame_emb = None
            self.frame_proj = None

        # UNet2DModel from diffusers - lightweight attention at deeper layers
        # Input channels: 1 (noisy target) + k (context images) + optional time
        in_ch = 1 + context_k + (1 if use_time_map else 0)
        self.unet = UNet2DModel(
            sample_size=IMG_SIZE
            , in_channels=in_ch
            , out_channels=1
            , layers_per_block=2
            , block_out_channels=(16, 32, 64, 128)
            , down_block_types=("DownBlock2D", "DownBlock2D", "DownBlock2D",
            "DownBlock2D")
            , up_block_types=("UpBlock2D", "UpBlock2D", "UpBlock2D",
            "UpBlock2D")
            , norm_num_groups = 16
        )

        def build_time_map(self, frame_idx: torch.Tensor | int, h: int, w: int) -> torch.Tensor:
            """Convert frame indices into a spatial map [B,1,H,W] for conditioning.
            """
            if not self.use_time_map:
                raise ValueError("build_time_map called but USE_TIME_MAP=False")
```

```

        if not torch.is_tensor(frame_idx):
            frame_idx = torch.tensor(frame_idx, device=next(self.parameters()).
↪device, dtype=torch.long)
        if frame_idx.dim() == 0:
            frame_idx = frame_idx.view(1)
        frame_idx = frame_idx.clamp(0, MAX_FRAME_INDEX).long()
        emb = self.frame_emb(frame_idx) # [B, TIME_EMB_DIM]
        scalar = self.frame_proj(emb) # [B, 1]
        return scalar.view(-1, 1, 1, 1).expand(-1, 1, h, w)

    def build_input(self, x_noisy: torch.Tensor, x_ctx: torch.Tensor, time_map:
↪torch.Tensor | None = None) -> torch.Tensor:
        """Assemble model input: [x_noisy, x_ctx, (optional) time_map]."""
        if self.use_time_map:
            if time_map is None:
                raise ValueError("time_map is required when USE_TIME_MAP=True")
            return torch.cat([x_noisy, x_ctx, time_map], dim=1)
        return torch.cat([x_noisy, x_ctx], dim=1)

    def forward(self, x_in, t):
        # x_in: [B, 1 + k (+1), H, W], t: [B] timesteps
        noise = self.unet(x_in, t).sample # [B, 1, H, W]
        return noise

    def q_sample(self, x0, t, noise=None):
        """Add noise to x0 at timestep t."""
        if noise is None:
            noise = torch.randn_like(x0)
        sqrt_alpha = sqrt_alphas_cumprod[t][:, None, None, None]
        sqrt_one_minus = sqrt_one_minus_alphas_cumprod[t][:, None, None, None]
        return sqrt_alpha * x0 + sqrt_one_minus * noise, noise

    def p_losses(self, x_tgt, x_ctx, time_map: torch.Tensor | None = None):
        """Training loss: predict the noise used to corrupt x_tgt."""
        B = x_tgt.size(0)
        t = torch.randint(0, Timesteps, (B,), device=x_tgt.device, dtype=torch.
↪long)
        x_noisy, noise = self.q_sample(x_tgt, t)
        x_in = self.build_input(x_noisy, x_ctx, time_map=time_map)
        pred = self.forward(x_in, t)
        mse = F.mse_loss(pred, noise)
        sqrt_alpha = sqrt_alphas_cumprod[t][:, None, None, None]
        sqrt_one_minus = sqrt_one_minus_alphas_cumprod[t][:, None, None, None]
        x0_pred = (x_noisy - sqrt_one_minus * pred) / (sqrt_alpha + 1e-8)
        return mse, x0_pred

```

1.9 8. Image preprocessing and dataset

We center-crop and resize frames, convert to grayscale tensors in $[-1, 1]$, and build **context**→**target** pairs (strictly adjacent in time).

```
[9]: def _preprocess_pil(img: Image.Image, size: int = IMG_SIZE) -> Image.Image:
    """Center-crop to square and resize to model input size."""
    w, h = img.size
    m = min(w, h)
    left = (w - m) // 2
    top = (h - m) // 2
    img = img.crop((left, top, left + m, top + m))
    if size is not None and (img.size[0] != size or img.size[1] != size):
        img = img.resize((size, size), Image.BICUBIC)
    return img

def load_pil(path: str, size: int = IMG_SIZE) -> Image.Image:
    """Load grayscale image and preprocess."""
    img = Image.open(path).convert('L')
    return _preprocess_pil(img, size=size)

def pil_to_tensor(img: Image.Image) -> torch.Tensor:
    """Convert PIL to float tensor in [-1, 1]."""
    arr = np.array(img, dtype=np.float32) / 255.0
    arr = arr * 2 - 1 # [-1, 1]
    return torch.from_numpy(arr).unsqueeze(0)

def _augment_stack(ctx_list: list[Image.Image], tgt: Image.Image):
    """Apply identical random augmentation to every context frame + target."""
    if not AUGMENT:
        return ctx_list, tgt

    do_hflip = AUGMENT_HFLIP > 0 and random.random() < AUGMENT_HFLIP
    do_vflip = AUGMENT_VFLIP > 0 and random.random() < AUGMENT_VFLIP
    angle = random.uniform(-AUGMENT_MAX_ROT, AUGMENT_MAX_ROT) if
    ↪AUGMENT_MAX_ROT and AUGMENT_MAX_ROT > 0 else 0.0
    brightness = 1.0 + random.uniform(-AUGMENT_BRIGHTNESS, AUGMENT_BRIGHTNESS)
    ↪if AUGMENT_BRIGHTNESS and AUGMENT_BRIGHTNESS > 0 else None
    contrast = 1.0 + random.uniform(-AUGMENT_CONTRAST, AUGMENT_CONTRAST) if
    ↪AUGMENT_CONTRAST and AUGMENT_CONTRAST > 0 else None

    def _apply(img: Image.Image) -> Image.Image:
        if do_hflip:
            img = TF.hflip(img)
        if do_vflip:
            img = TF.vflip(img)
        if angle != 0.0:
```

```

        img = TF.rotate(img, angle, interpolation=TF.InterpolationMode.
↪BILINEAR)
        if brightness is not None:
            img = TF.adjust_brightness(img, brightness)
        if contrast is not None:
            img = TF.adjust_contrast(img, contrast)
        return img

    ctx_list = [_apply(img) for img in ctx_list]
    tgt = _apply(tgt)
    return ctx_list, tgt

def load_image(path: str, size: int = IMG_SIZE):
    """Load image and return tensor [1,H,W] in [-1, 1]."""
    img = load_pil(path, size=size)
    return pil_to_tensor(img)

def denorm(x):
    """Convert [-1,1] -> [0,1] for visualization."""
    return (x.clamp(-1,1) + 1) / 2

```

1.9.1 8.1 Dataset (context → target windows)

Builds deterministic context windows and caches transition indices for analysis/metrics.

```

[10]: class PairDataset(Dataset):
    """Sharma-style windowed next-frame pairs (strict adjacency).
    Context frames [t-K+1..t] predict target frame t+1 (delta_t = 1).
    """
    def __init__(
        self,
        index_df: pd.DataFrame,
        embryo_ids: set[str],
        split_seed: int,
        split_name: str,
        build_transition_index: bool = True,
    ):
        # Keep only rows that belong to the split's embryo IDs
        self.df = index_df[index_df['embryo_id'].isin(set(embryo_ids))].copy()
↪reset_index(drop=True)
        self.split_name = split_name
        self.build_transition_index = build_transition_index
        self.rng = np.random.RandomState(split_seed + {'train': 0, 'val': 1,
↪'test': 2}.get(split_name, 0))
        self.sequences: list[tuple[str, str]] = []
        self.sequences_by_plane: dict[str, list[tuple[str, str]]] = {}
        self.frames_by_seq: dict[tuple[str, str], list[int]] = {}

```

```

self.path_by_seq: dict[tuple[str, str], dict[int, str]] = {}
self.phase_by_seq: dict[tuple[str, str], dict[int, int]] = {}
self.valid_t_ends_by_seq: dict[tuple[str, str], list[int]] = {}
self.transition_t_ends_by_seq: dict[tuple[str, str], list[int]] = {}
self.non_transition_t_ends_by_seq: dict[tuple[str, str], list[int]] = {}
self.max_frame_by_seq: dict[tuple[str, str], int] = {}
self.windows: list[dict] = []
self.windows_by_plane: dict[str, list[dict]] = {}
self.transition_windows_by_plane: dict[str, list[dict]] = {}
self.non_transition_windows_by_plane: dict[str, list[dict]] = {}
self._skipped_short = 0
self._skipped_no_tgt = 0
self._build_index()

def _build_index(self):
    # Group by embryo + plane, then build all valid context windows
    for (eid, plane), g in self.df.groupby(['embryo_id', 'plane']):
        g = g.sort_values('frame_num').reset_index(drop=True)
        frame_nums = g['frame_num'].tolist()
        if len(frame_nums) <= CONTEXT_K:
            self._skipped_short += 1
            continue
        frame_set = set(frame_nums)
        path_by = dict(zip(frame_nums, g['path'].tolist()))
        phase_by = dict(zip(frame_nums, g['phase_id'].tolist()))
        max_frame = max(frame_nums)
        self.max_frame_by_seq[(eid, plane)] = max_frame

        valid_t_ends = []
        for t_end in frame_nums:
            # Require full context [t-K+1..t] and adjacent target t+1
            if (t_end + 1) not in frame_set:
                continue
            ok = True
            for f in range(t_end - CONTEXT_K + 1, t_end + 1):
                if f not in frame_set:
                    ok = False
                    break
            if ok:
                valid_t_ends.append(t_end)
        if not valid_t_ends:
            self._skipped_no_tgt += 1
            continue

        if self.build_transition_index:
            transition_t_ends = []
            non_transition_t_ends = []

```



```

        for t in valid_t_ends:
            phase_t = phase_by.get(t)
            phase_next = phase_by.get(t + 1)
            if phase_next is None:
                continue
            # Transition windows: target frame t+1 is first frame of
↳new phase
            if phase_next != phase_t:
                transition_t_ends.append(t)
            else:
                non_transition_t_ends.append(t)
        else:
            transition_t_ends = []
            non_transition_t_ends = []

    seq = (eid, plane)
    self.sequences.append(seq)
    self.sequences_by_plane.setdefault(plane, []).append(seq)
    self.frames_by_seq[seq] = frame_nums
    self.path_by_seq[seq] = path_by
    self.phase_by_seq[seq] = phase_by
    self.valid_t_ends_by_seq[seq] = valid_t_ends
    self.transition_t_ends_by_seq[seq] = transition_t_ends
    self.non_transition_t_ends_by_seq[seq] = non_transition_t_ends

    # Build windows (overlapping, deterministic)
    for t_end in valid_t_ends:
        row = {
            'embryo_id': eid,
            'plane': plane,
            't_end': t_end,
            'tgt_frame': t_end + 1,
            'max_frame': max_frame,
        }
        self.windows.append(row)
        self.windows_by_plane.setdefault(plane, []).append(row)
        if self.build_transition_index and t_end in transition_t_ends:
            self.transition_windows_by_plane.setdefault(plane, []).
↳append(row)
        elif self.build_transition_index:
            self.non_transition_windows_by_plane.setdefault(plane, []).
↳append(row)

    if VERBOSE:
        frames_per_seq = [len(self.frames_by_seq[s]) for s in self.
↳sequences]

```

```

        valid_tgts_per_seq = [len(self.valid_t_ends_by_seq[s]) for s in
↪self.sequences]
        trans_tgts_per_seq = [len(self.transition_t_ends_by_seq[s]) for s
↪in self.sequences]
        total_frames = int(np.sum(frames_per_seq)) if frames_per_seq else 0
        def _avg(vals):
            return float(np.mean(vals)) if vals else 0.0
        def _minmax(vals):
            return (int(np.min(vals)), int(np.max(vals))) if vals else (0,
↪0)

        fp_min, fp_max = _minmax(frames_per_seq)
        vt_min, vt_max = _minmax(valid_tgts_per_seq)
        tr_min, tr_max = _minmax(trans_tgts_per_seq)
        print(
            f"[PairDataset:{self.split_name}] sequences={len(self.
↪sequences)} | "
            f"frames={total_frames} | frames/seq avg={_avg(frames_per_seq):.
↪1f} (min={fp_min}, max={fp_max}) | "
            f"valid windows/seq avg={_avg(valid_tgts_per_seq):.1f}
↪(min={vt_min}, max={vt_max}) | "
            f"transition windows/seq avg={_avg(trans_tgts_per_seq):.1f}
↪(min={tr_min}, max={tr_max}) | "
            f"skipped short={self._skipped_short}, skipped no-target={self.
↪_skipped_no_tgt}"
        )

    def __len__(self):
        return len(self.windows)

    def _build_row(self, row: dict):
        # Expand a window row into actual file paths
        seq = (row['embryo_id'], row['plane'])
        t_end = row['t_end']
        ctx_frames = [t_end - i for i in range(CONTEXT_K - 1, -1, -1)]
        ctx_paths = [self.path_by_seq[seq][f] for f in ctx_frames]
        tgt_path = self.path_by_seq[seq][t_end + 1]
        out = dict(row)
        out['ctx_paths'] = ctx_paths
        out['tgt_path'] = tgt_path
        return out

    def sample_pairs_for_plane(self, plane: str, n: int, mode: str = 'any'):
        """Utility for visual sampling: transitions vs non-transitions vs all.
↪"""
        if mode == 'transition':
            rows = self.transition_windows_by_plane.get(plane, [])

```

```

elif mode == 'non_transition':
    rows = self.non_transition_windows_by_plane.get(plane, [])
else:
    rows = self.windows_by_plane.get(plane, [])
if not rows or n <= 0:
    return []
picks = [rows[self.rng.randint(len(rows))] for _ in range(n)]
return [self._build_row(r) for r in picks]

def sample_transition_pairs_for_plane(self, plane: str, n: int):
    return self.sample_pairs_for_plane(plane, n, mode='transition')

def sample_non_transition_pairs_for_plane(self, plane: str, n: int):
    return self.sample_pairs_for_plane(plane, n, mode='non_transition')

def __getitem__(self, idx):
    # Returns tensors for one training example
    row = self._build_row(self.windows[idx])
    ctx_pils = [load_pil(p) for p in row['ctx_paths']]
    tgt_pil = load_pil(row['tgt_path'])
    ctx_pils, tgt_pil = _augment_stack(ctx_pils, tgt_pil)
    x_ctx = torch.stack([pil_to_tensor(p).squeeze(0) for p in ctx_pils],
    ↪dim=0) # [k, H, W]
    x_tgt = pil_to_tensor(tgt_pil) # [1, H, W]
    # Absolute frame index embedding (target frame)
    if USE_TIME_MAP:
        frame_idx = min(int(row['tgt_frame']), int(MAX_FRAME_INDEX))
        frame_idx = torch.tensor(frame_idx, dtype=torch.long)
        return x_ctx, x_tgt, frame_idx
    return x_ctx, x_tgt

```

1.10 9. DataLoaders

We build datasets once and sample training pairs on the fly. Validation/test datasets also keep a **transition index** for specialized metrics.

```

[11]: # Build datasets/loaders once; pairs are sampled on the fly
train_dataset = PairDataset(index_for_pairs, train_ids, SPLIT_SEED, 'train',
    ↪build_transition_index=False)
val_dataset = PairDataset(index_for_pairs, val_ids, SPLIT_SEED, 'val',
    ↪build_transition_index=True)
test_dataset = PairDataset(index_for_pairs, test_ids, SPLIT_SEED, 'test',
    ↪build_transition_index=True)

train_loader = DataLoader(train_dataset, batch_size=BATCH_SIZE, shuffle=True,
    ↪num_workers=NUM_WORKERS, pin_memory=True, persistent_workers=True)

```

```
val_loader = DataLoader(val_dataset, batch_size=BATCH_SIZE, shuffle=False,
    ↪num_workers=NUM_WORKERS, pin_memory=True)
test_loader = DataLoader(test_dataset, batch_size=BATCH_SIZE, shuffle=False,
    ↪num_workers=NUM_WORKERS, pin_memory=True)
```

```
[PairDataset:train] sequences=565 | frames=230158 | frames/seq avg=407.4
(min=106, max=647) | valid windows/seq avg=398.2 (min=99, max=640) | transition
windows/seq avg=0.0 (min=0, max=0) | skipped short=0, skipped no-target=0
[PairDataset:val] sequences=69 | frames=30890 | frames/seq avg=447.7 (min=179,
max=639) | valid windows/seq avg=437.9 (min=165, max=632) | transition
windows/seq avg=11.7 (min=6, max=14) | skipped short=0, skipped no-target=0
[PairDataset:test] sequences=69 | frames=28627 | frames/seq avg=414.9 (min=145,
max=564) | valid windows/seq avg=405.6 (min=138, max=554) | transition
windows/seq avg=10.3 (min=4, max=14) | skipped short=0, skipped no-target=0
```

1.11 10. Sampling helpers and guidance

This section defines helper utilities for **DDIM sampling**, optional **external classifier guidance**, and **DDIM inversion** (used in the inversion regularizer).

```
[12]: diffusion = DiffusionCond(context_k=CONTEXT_K, use_time_map=USE_TIME_MAP).
    ↪to(device)
if USE_GRAD_CHECKPOINTING:
    diffusion.unet.enable_gradient_checkpointing()
optimizer = torch.optim.AdamW(diffusion.parameters(), lr=LR, betas=(0.9, 0.999))
```

1.11.1 10.1 Sampling & guidance helpers

Utilities for classifier-free sampling, optional external classifier guidance, and DDIM inversion.

```
[13]: def _repeat_channels(x, target_channels=3):
    """Repeat single-channel tensors along channel dim."""
    if x.shape[1] == target_channels:
        return x
    return x.repeat(1, target_channels, 1, 1)

def _eps_from_model(x_noisy: torch.Tensor, x_ctx: torch.Tensor, t_tensor: torch.
    ↪Tensor, time_map: torch.Tensor | None = None):
    """Query UNet to predict noise epsilon."""
    x_in = diffusion.build_input(x_noisy, x_ctx, time_map=time_map)
    return diffusion.forward(x_in, t_tensor)

def _ensure_context_batch(x_ctx: torch.Tensor | None, shape):
    """Ensure context tensor is [B, k, H, W]."""
    B, H, W = shape[0], shape[2], shape[3]
    if x_ctx is None:
        return torch.zeros((B, CONTEXT_K, H, W), device=device)
    x_ctx = x_ctx.to(device)
```

```

    if x_ctx.dim() == 3:
        x_ctx = x_ctx.unsqueeze(0)
    if x_ctx.shape[1] != CONTEXT_K:
        raise ValueError(f"Expected context with k={CONTEXT_K} channels, got {x_ctx.shape[1]}")
    return x_ctx

def _ensure_time_map(frame_idx: torch.Tensor | int | None, shape):
    """Build time map if enabled; otherwise return None."""
    if not USE_TIME_MAP:
        return None
    if frame_idx is None:
        raise ValueError("frame_idx is required when USE_TIME_MAP=True")
    if not torch.is_tensor(frame_idx):
        frame_idx = torch.tensor(frame_idx, device=device, dtype=torch.long)
    if frame_idx.dim() == 0:
        frame_idx = frame_idx.view(1)
    frame_idx = frame_idx.clamp(0, MAX_FRAME_INDEX).long()
    return diffusion.build_time_map(frame_idx, shape[2], shape[3])

def ddim_sample_from_model(
    shape, x_ctx, frame_idx=None,
    steps=DDIM_STEPS, eta=DDIM_ETA,
):
    """DDIM sampling."""
    B = shape[0]
    x = torch.randn((B, 1, shape[2], shape[3]), device=device)
    x_ctx = _ensure_context_batch(x_ctx, shape)
    time_map = _ensure_time_map(frame_idx, shape)
    ddim_scheduler.set_timesteps(steps, device=device)
    for t in ddim_scheduler.timesteps:
        t_tensor = torch.full((B,), int(t), device=device, dtype=torch.long)
        with torch.no_grad():
            eps = _eps_from_model(x, x_ctx, t_tensor, time_map=time_map)
            x = ddim_scheduler.step(eps, t, x, eta=eta).prev_sample
    return x.detach()

def ddim_invert(x0, x_ctx, frame_idx=None, steps=INVERSION_STEPS):
    """DDIM inversion: map  $x_0 \rightarrow x_T$  deterministically ( $\eta=0$ )."""
    B = x0.size(0)
    x_ctx = _ensure_context_batch(x_ctx, (B, 1, x0.shape[-2], x0.shape[-1]))
    time_map = _ensure_time_map(frame_idx, (B, 1, x0.shape[-2], x0.shape[-1]))
    x = x0
    ddim_scheduler.set_timesteps(steps, device=device)
    timesteps = list(reversed(ddim_scheduler.timesteps))
    alphas_cumprod = ddim_scheduler.alphas_cumprod.to(device)
    for i, t in enumerate(timesteps[:-1]):

```

```

        t_next = timesteps[i + 1]
        t_tensor = torch.full((B,), int(t), device=device, dtype=torch.long)
        with torch.no_grad():
            eps = _eps_from_model(x, x_ctx, t_tensor, time_map=time_map)
            alpha_t = alphas_cumprod[int(t)].view(1, 1, 1, 1)
            alpha_next = alphas_cumprod[int(t_next)].view(1, 1, 1, 1)
            sqrt_alpha_t = torch.sqrt(alpha_t)
            sqrt_one_minus_alpha_t = torch.sqrt(1.0 - alpha_t)
            x0_pred = (x - sqrt_one_minus_alpha_t * eps) / sqrt_alpha_t
            x = torch.sqrt(alpha_next) * x0_pred + torch.sqrt(1.0 - alpha_next) * ␣
↪eps
        return x.detach()

def ddim_reconstruct_from_xT(
    x_T, x_ctx, frame_idx=None,
    steps=INVERSION_STEPS,
):
    """DDIM denoising from x_T back to x0 (supports gradients)."""
    B = x_T.size(0)
    x_ctx = _ensure_context_batch(x_ctx, (B, 1, x_T.shape[-2], x_T.shape[-1]))
    time_map = _ensure_time_map(frame_idx, (B, 1, x_T.shape[-2], x_T.shape[-1]))
    x = x_T
    ddim_scheduler.set_timesteps(steps, device=device)
    for t in ddim_scheduler.timesteps:
        t_tensor = torch.full((B,), int(t), device=device, dtype=torch.long)
        eps = _eps_from_model(x, x_ctx, t_tensor, time_map=time_map)
        x = ddim_scheduler.step(eps, t, x, eta=0.0).prev_sample
    return x

def sample_from_model(shape, x_ctx, frame_idx=None):
    """Generate samples using DDIM."""
    return ddim_sample_from_model(
        shape, x_ctx, frame_idx=frame_idx,
        steps=DDIM_STEPS, eta=DDIM_ETA
    )

```

```

[14]: def save_tensor_image(tensor, path):
    """Save tensor as grayscale image."""
    arr = (denorm(tensor).clamp(0, 1).cpu().numpy() * 255).astype(np.uint8)
    if arr.ndim == 3 and arr.shape[0] == 1:
        arr = arr[0]
    Image.fromarray(arr).save(path)

lpips_metric = None
if COMPUTE_VAL_METRICS or (LPIPS_LOSS_WEIGHT and LPIPS_LOSS_WEIGHT > 0):
    lpips_metric = LPIPS(net='alex').to(device)
    lpips_metric.eval()

```

```

    for p in lpips_metric.parameters():
        p.requires_grad_(False)

_gaussian_kernel_cache: dict[tuple[int, float, str, int | None, torch.dtype],
    ↪torch.Tensor] = {}
def _gaussian_kernel(window_size: int, sigma: float, device: torch.device,
    ↪dtype: torch.dtype):
    """Cacheable 2D Gaussian window for SSIM."""
    key = (window_size, sigma, device.type, device.index, dtype)
    kernel = _gaussian_kernel_cache.get(key)
    if kernel is None:
        coords = torch.arange(window_size, device=device, dtype=dtype) -
    ↪window_size // 2
        gauss = torch.exp(-(coords ** 2) / (2 * sigma ** 2))
        gauss /= gauss.sum()
        kernel_1d = gauss.unsqueeze(0)
        kernel_2d = kernel_1d.t() @ kernel_1d
        kernel_2d /= kernel_2d.sum()
        kernel = kernel_2d.unsqueeze(0).unsqueeze(0)
        _gaussian_kernel_cache[key] = kernel
    return kernel

def compute_ssim_metric(pred: torch.Tensor, target: torch.Tensor, window_size:
    ↪int = 11, sigma: float = 1.5) -> torch.Tensor:
    """Compute SSIM in [-1,1] space by internally mapping to [0,1]."""
    pred = pred.clamp(-1, 1)
    target = target.clamp(-1, 1)
    pred_01 = (pred + 1) / 2
    target_01 = (target + 1) / 2
    pred_01 = pred_01.float()
    target_01 = target_01.float()
    device = pred_01.device
    dtype = pred_01.dtype
    channels = pred_01.size(1)
    kernel = _gaussian_kernel(window_size, sigma, device, dtype).
    ↪to(device=device, dtype=dtype)
    kernel = kernel.expand(channels, 1, window_size, window_size)
    padding = window_size // 2
    mu_pred = F.conv2d(pred_01, kernel, padding=padding, groups=channels)
    mu_target = F.conv2d(target_01, kernel, padding=padding, groups=channels)
    mu_pred_sq = mu_pred.pow(2)
    mu_target_sq = mu_target.pow(2)
    mu_pred_target = mu_pred * mu_target
    sigma_pred = F.conv2d(pred_01 * pred_01, kernel, padding=padding,
    ↪groups=channels) - mu_pred_sq

```

```

    sigma_target = F.conv2d(target_01 * target_01, kernel, padding=padding,
↪groups=channels) - mu_target_sq
    sigma_pred_target = F.conv2d(pred_01 * target_01, kernel, padding=padding,
↪groups=channels) - mu_pred_target
    sigma_pred = torch.clamp_min(sigma_pred, 0.0)
    sigma_target = torch.clamp_min(sigma_target, 0.0)
    c1 = (0.01) ** 2
    c2 = (0.03) ** 2
    numerator = (2 * mu_pred_target + c1) * (2 * sigma_pred_target + c2)
    denominator = (mu_pred_sq + mu_target_sq + c1) * (sigma_pred + sigma_target
↪+ c2)
    ssim_map = numerator / (denominator + 1e-6)
    return ssim_map.mean()

def compute_lpips_metric(pred: torch.Tensor, target: torch.Tensor) -> torch.
↪Tensor:
    """Compute LPIPS on grayscale by repeating to 3 channels."""
    if lpips_metric is None:
        raise RuntimeError('LPIPS metric requested but the LPIPS model is not
↪initialised.')
    pred_in = _repeat_channels(pred.clamp(-1, 1))
    target_in = _repeat_channels(target.clamp(-1, 1))
    with torch.no_grad():
        return lpips_metric(pred_in, target_in, normalize=False).mean()

```

Setting up [LPIPS] perceptual loss: trunk [alex], v[0.1], spatial [off]

```

/home/khasion/miniconda3/envs/thesis/lib/python3.10/site-
packages/torchvision/models/_utils.py:208: UserWarning: The parameter
'pretrained' is deprecated since 0.13 and may be removed in the future, please
use 'weights' instead.

```

```

warnings.warn(
/home/khasion/miniconda3/envs/thesis/lib/python3.10/site-
packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a
weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed
in the future. The current behavior is equivalent to passing
`weights=AlexNet_Weights.IMAGENET1K_V1`. You can also use
`weights=AlexNet_Weights.DEFAULT` to get the most up-to-date weights.
warnings.warn(msg)

```

```

Loading model from: /home/khasion/miniconda3/envs/thesis/lib/python3.10/site-
packages/lpips/weights/v0.1/alex.pth

```

1.12 12. Training loop

Training predicts noise on the target frame, optionally adds an **inversion regularizer**, computes metrics on validation pairs, and saves checkpoints + sample grids.


```

[15]: # Train loop - context-conditioned diffusion model
best_val = None
ckpt_path = OUT_MODELS_DIR / f'nextframe_ctxk{CONTEXT_K}_{IMG_SIZE}px.pt'

# Early stopping config
EARLY_STOPPING_PATIENCE = 15 # set to None or 0 to disable
epochs_no_improve = 0

# GradScaler for proper mixed precision training (prevents NaN from underflow/
↳ overflow)
use_amp = torch.cuda.is_available()
scaler = torch.amp.GradScaler('cuda', enabled=use_amp)
metric_history = []

# Learning-rate scheduler setup (ReduceLROnPlateau only)
lr_scheduler = None
if USE_LR_SCHEDULER:
    lr_scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(
        optimizer, mode='min', factor=0.8, patience=2, min_lr=LR_MIN
    )

def one_epoch(loader, train=True, desc: str | None = None):
    """Run one epoch of training or validation over a loader."""
    diffusion.train(mode=train)
    total = 0.0
    count = 0
    nan_batches = 0
    inv_total = 0.0
    inv_count = 0
    if desc is None:
        desc = "Train" if train else "Val"
    for batch_idx, batch in enumerate(tqdm_bar(loader, desc=desc)):
        if USE_TIME_MAP:
            x_ctx, x_tgt, frame_idx = batch
            frame_idx = frame_idx.to(device, non_blocking=True)
        else:
            x_ctx, x_tgt = batch
            frame_idx = None
        x_ctx = x_ctx.to(device, non_blocking=True) # [B, k, H, W]
        x_tgt = x_tgt.to(device, non_blocking=True) # [B, 1, H, W]

        time_map = diffusion.build_time_map(frame_idx, x_ctx.shape[-2], x_ctx.
↳ shape[-1]) if USE_TIME_MAP else None

        # Mixed precision training with proper GradScaler
        with torch.autocast(device_type='cuda', enabled=use_amp):

```

```

        mse_loss, x0_pred = diffusion.p_losses(x_tgt, x_ctx,
↪time_map=time_map)
        loss = mse_loss
        if X0_L1_WEIGHT and X0_L1_WEIGHT > 0:
            x0_l1 = F.l1_loss(x0_pred, x_tgt)
            loss = loss + X0_L1_WEIGHT * x0_l1
        if X0_SSIM_WEIGHT and X0_SSIM_WEIGHT > 0:
            ssim_val = compute_ssim_metric(x0_pred, x_tgt)
            loss = loss + X0_SSIM_WEIGHT * (1.0 - ssim_val)
        if lpips_metric is not None:
            # Optional perceptual loss computed on predicted x0
            lpips_in = _repeat_channels(x0_pred.clamp(-1, 1))
            tgt_in = _repeat_channels(x_tgt.clamp(-1, 1))
            lpips_loss = lpips_metric(lpips_in, tgt_in, normalize=False).
↪mean()

            loss = loss + LPIPS_LOSS_WEIGHT * lpips_loss

        # Optional inversion training (DDIM inversion + reconstruction)
        inv_loss = None
        if INVERSION_TRAINING and train and (batch_idx % max(1,
↪INVERSION_EVERY_N) == 0):
            b_inv = min(x_tgt.size(0), INVERSION_MAX_BATCH)
            x_ctx_inv = x_ctx[:b_inv]
            x_tgt_inv = x_tgt[:b_inv]
            frame_idx_inv = frame_idx[:b_inv] if frame_idx is not None else None
            with torch.no_grad():
                x_T = ddim_invert(
                    x_tgt_inv, x_ctx_inv, frame_idx=frame_idx_inv,
                    steps=INVERSION_STEPS
                )
            with torch.autocast(device_type='cuda', enabled=use_amp):
                x_recon = ddim_reconstruct_from_xT(
                    x_T, x_ctx_inv, frame_idx=frame_idx_inv,
                    steps=INVERSION_STEPS
                )
            inv_loss = F.l1_loss(x_recon, x_tgt_inv)
            loss = loss + INVERSION_LAMBDA * inv_loss

        # Check for NaN/Inf BEFORE backward (critical: don't corrupt model with
↪bad gradients)
        if torch.isnan(loss) or torch.isinf(loss):
            nan_batches += 1
            if nan_batches <= 3: # Only warn first few times
                print(f"WARNING: NaN/Inf loss detected, skipping batch (total:
↪{nan_batches})")
            continue

```

```

    if train:
        optimizer.zero_grad(set_to_none=True)
        scaler.scale(loss).backward()
        scaler.unscale_(optimizer)
        torch.nn.utils.clip_grad_norm_(diffusion.parameters(), GRAD_CLIP)
        scaler.step(optimizer)
        scaler.update()

    total += loss.item() * x_tgt.size(0)
    count += x_tgt.size(0)
    if inv_loss is not None:
        inv_total += inv_loss.item() * x_tgt.size(0)
        inv_count += x_tgt.size(0)

    if nan_batches > 0:
        print(f" (Epoch had {nan_batches} NaN/Inf batches skipped)")
    avg = total / max(1, count)
    inv_avg = inv_total / max(1, inv_count) if inv_count > 0 else None
    return avg, inv_avg

def _rows_to_tensors(rows: list[dict]):
    """Turn sampled metadata rows into batched tensors."""
    ctx_list = []
    tgt_list = []
    frame_idx_list = []
    for row in rows:
        ctx_tensors = [load_image(p).squeeze(0) for p in row['ctx_paths']] # [k, H, W]
        x_ctx = torch.stack(ctx_tensors, dim=0)
        x_tgt = load_image(row['tgt_path']) # [1, H, W]
        ctx_list.append(x_ctx)
        tgt_list.append(x_tgt)
        if USE_TIME_MAP:
            frame_idx = min(int(row['tgt_frame']), int(MAX_FRAME_INDEX))
            frame_idx_list.append(frame_idx)
    x_ctx = torch.stack(ctx_list, dim=0)
    x_tgt = torch.stack(tgt_list, dim=0)
    if USE_TIME_MAP:
        frame_idx = torch.tensor(frame_idx_list, dtype=torch.long)
        return x_ctx, x_tgt, frame_idx
    return x_ctx, x_tgt

def _compute_metrics_from_rows(rows: list[dict], desc: str = "Val metrics"):
    """Compute SSIM/LPIPS for a list of sampled pairs."""
    if not rows:
        return None, None
    ssim_vals = []

```

```

lpips_vals = []
for i in range(0, len(rows), BATCH_SIZE):
    batch_rows = rows[i:i + BATCH_SIZE]
    if USE_TIME_MAP:
        x_ctx, x_tgt, frame_idx = _rows_to_tensors(batch_rows)
        frame_idx = frame_idx.to(device, non_blocking=True)
    else:
        x_ctx, x_tgt = _rows_to_tensors(batch_rows)
        frame_idx = None
    x_ctx = x_ctx.to(device, non_blocking=True)
    x_tgt = x_tgt.to(device, non_blocking=True)
    x_gen = sample_from_model(
        (x_ctx.size(0), 1, IMG_SIZE, IMG_SIZE),
        x_ctx=x_ctx,
        frame_idx=frame_idx,
    )
    ssim_vals.append(compute_ssim_metric(x_gen, x_tgt).item())
    if lpips_metric is not None:
        lpips_vals.append(compute_lpips_metric(x_gen, x_tgt).item())
ssim = float(np.mean(ssim_vals)) if ssim_vals else None
lpips = float(np.mean(lpips_vals)) if lpips_vals else None
return ssim, lpips

def _collect_pairs_for_metrics(dataset: PairDataset, mode: str, max_samples:
    ↪int):
    """Sample rows across planes for metric evaluation."""
    rows = []
    planes = list(dataset.windows_by_plane.keys())
    for plane in planes:
        if len(rows) >= max_samples:
            break
        n_needed = max_samples - len(rows)
        if mode == 'transition':
            rows.extend(dataset.sample_transition_pairs_for_plane(plane,
                ↪n_needed))
        elif mode == 'non_transition':
            rows.extend(dataset.sample_non_transition_pairs_for_plane(plane,
                ↪n_needed))
        else:
            rows.extend(dataset.sample_pairs_for_plane(plane, n_needed))
    return rows

def _save_sample_grid(rows: list[dict], out_dir: Path, title_suffix: str):
    """Save context/target/generated grids for qualitative inspection."""
    if not rows:
        return
    out_dir.mkdir(parents=True, exist_ok=True)

```

```

for idx, row in enumerate(rows):
    ctx_tensors = [load_image(p).squeeze(0) for p in row['ctx_paths']]
    tgt_tensor_cpu = load_image(row['tgt_path'])
    x_ctx = torch.stack(ctx_tensors, dim=0).unsqueeze(0).to(device)
    frame_idx = None
    if USE_TIME_MAP:
        frame_idx = min(int(row['tgt_frame']), int(MAX_FRAME_INDEX))
        frame_idx = torch.tensor([frame_idx], device=device, dtype=torch.
↳ long)

    x_gen = sample_from_model(
        (1, 1, IMG_SIZE, IMG_SIZE),
        x_ctx=x_ctx,
        frame_idx=frame_idx,
    )
    # number of panels: k context + target + generated
    ncols = CONTEXT_K + 2
    fig, axes = plt.subplots(1, ncols, figsize=(3 * ncols, 3))
    if ncols == 1:
        axes = [axes]

    visuals = []
    titles = []

    # Context frames: ctx_tensors[i] are [H,W] in [-1,1]
    for i, ctx in enumerate(ctx_tensors):
        # i=0 corresponds to t-k, i=k-1 corresponds to t-1
        img = denorm(ctx.unsqueeze(0).unsqueeze(0))[0, 0].cpu().numpy()
        visuals.append(img)
        titles.append(f"Ctx (t-{{CONTEXT_K - i}})")

    # Target frame: tgt_tensor_cpu is [1,H,W] in [-1,1]
    tgt_img = denorm(tgt_tensor_cpu.unsqueeze(0))[0, 0].cpu().numpy()
    visuals.append(tgt_img)
    titles.append("Target (t+1)")

    # Generated frame: x_gen is [1,1,H,W] in [-1,1]
    gen_img = denorm(x_gen)[0, 0].detach().cpu().numpy()
    visuals.append(gen_img)
    titles.append("Generated")

    for ax, img, title in zip(axes, visuals, titles):
        ax.imshow(img, cmap="gray")
        ax.set_title(title)
        ax.axis("off")

    fig.suptitle(f"{{row['embryo_id']}} ({{row['plane']}}) - context_
↳ k={{CONTEXT_K}} → next {{title_suffix}}", fontsize=10)

```

```

fig.tight_layout()
save_path = out_dir / f'sample_{idx:02d}.png'
fig.savefig(save_path, dpi=150)
plt.close(fig)

def _save_rollout_grid(row: dict, out_dir: Path, steps: int):
    """Roll out multiple steps by feeding generated frames back into context."""
    out_dir.mkdir(parents=True, exist_ok=True)
    ctx_tensors = [load_image(p).squeeze(0) for p in row['ctx_paths']]
    x_ctx = torch.stack(ctx_tensors, dim=0).unsqueeze(0).to(device)
    t_end = int(row['t_end'])

    generated = []
    frame_idx = min(int(t_end + 1), int(MAX_FRAME_INDEX))
    for s in range(steps):
        if USE_TIME_MAP:
            frame_idx = min(frame_idx, int(MAX_FRAME_INDEX))
            frame_idx_tensor = torch.tensor([frame_idx], device=device,
dtype=torch.long)
        else:
            frame_idx_tensor = None
        x_gen = sample_from_model(
            (1, 1, IMG_SIZE, IMG_SIZE),
            x_ctx=x_ctx,
            frame_idx=frame_idx_tensor,
        )
        generated.append(x_gen[0].detach().cpu())
        # update context window: drop oldest, append generated
        x_ctx = torch.cat([x_ctx[:, 1:], x_gen.detach()], dim=1)
        frame_idx += 1

    ncols = CONTEXT_K + steps
    fig, axes = plt.subplots(1, ncols, figsize=(3 * ncols, 3))
    if ncols == 1:
        axes = [axes]

    # Context frames
    for i, ctx in enumerate(ctx_tensors):
        img = denorm(ctx.unsqueeze(0))[0].cpu().numpy()
        axes[i].imshow(img, cmap="gray")
        axes[i].set_title(f"Ctx (t-{CONTEXT_K - i})")
        axes[i].axis("off")

    # Generated rollout frames
    for i, gen in enumerate(generated):
        img = denorm(gen)[0].cpu().numpy()
        axes[CONTEXT_K + i].imshow(img, cmap="gray")

```

```

        axes[CONTEXT_K + i].set_title(f"Gen (t+{i+1})")
        axes[CONTEXT_K + i].axis("off")

    fig.suptitle(f"{row['embryo_id']} ({row['plane']}) - rollout {steps}_
↳steps", fontsize=10)
    fig.tight_layout()
    save_path = out_dir / f"rollout_t{row['t_end']:05d}_k{CONTEXT_K}_h{steps}.
↳png"
    fig.savefig(save_path, dpi=150)
    plt.close(fig)

```

1.12.1 12.1 Run training

Main loop for training, validation, checkpointing, and sample generation.

```

[ ]: for epoch in range(1, EPOCHS+1):
    tl, tl_inv = one_epoch(train_loader, train=True, desc=f"Epoch {epoch:03d}_
↳Train")
    with torch.no_grad():
        vl, _ = one_epoch(val_loader, train=False, desc=f"Epoch {epoch:03d}_
↳Val")

    if lr_scheduler is not None:
        lr_scheduler.step(vl)

    val_ssim = None
    val_lpips = None
    val_ssim_trans = None
    val_lpips_trans = None
    val_ssim_non = None
    val_lpips_non = None
    if COMPUTE_VAL_METRICS and VAL_METRIC_SAMPLES > 0:
        diffusion.eval()

        # Overall metrics (mixed transition + non-transition)
        all_rows = _collect_pairs_for_metrics(val_dataset, mode='any',
↳max_samples=VAL_METRIC_SAMPLES)
        val_ssim, val_lpips = _compute_metrics_from_rows(all_rows, desc="Val_
↳metrics (all)")

        # Transitional-only metrics
        trans_rows = _collect_pairs_for_metrics(val_dataset, mode='transition',
↳max_samples=VAL_METRIC_SAMPLES)
        val_ssim_trans, val_lpips_trans = _
↳compute_metrics_from_rows(trans_rows, desc="Val metrics (transition)")

        # Non-transitional metrics

```

```

        non_rows = _collect_pairs_for_metrics(val_dataset,
↪mode='non_transition', max_samples=VAL_METRIC_SAMPLES)
        val_ssim_non, val_lpips_non = _compute_metrics_from_rows(non_rows,
↪desc="Val metrics (non-transition)")

    log_msg = f"[Epoch {epoch:03d}] train {tl:.4f} | val {vl:.4f}"
    if tl_inv is not None:
        log_msg += f" | inv {tl_inv:.4f}"
    if val_ssim is not None:
        log_msg += f" | SSIM {val_ssim:.4f}"
    if val_ssim_trans is not None:
        log_msg += f" | SSIM_tr {val_ssim_trans:.4f}"
    if val_ssim_non is not None:
        log_msg += f" | SSIM_non {val_ssim_non:.4f}"
    if val_lpips is not None:
        log_msg += f" | LPIPS {val_lpips:.4f}"
    if val_lpips_trans is not None:
        log_msg += f" | LPIPS_tr {val_lpips_trans:.4f}"
    if val_lpips_non is not None:
        log_msg += f" | LPIPS_non {val_lpips_non:.4f}"
    print(log_msg)
    metric_history.append({
        'train_loss': tl,
        'val_loss': vl,
        'train_inv': tl_inv,
        'val_ssim': val_ssim,
        'val_ssim_tr': val_ssim_trans,
        'val_ssim_non': val_ssim_non,
        'val_lpips': val_lpips,
        'val_lpips_tr': val_lpips_trans,
        'val_lpips_non': val_lpips_non
    })

    if best_val is None or vl < best_val - 1e-8:
        best_val = vl
        epochs_no_improve = 0
        torch.save({'model': diffusion.state_dict(), 'epoch': epoch, 'val_loss':
↪ vl, 'config': {
            'IMG_SIZE': IMG_SIZE, 'Timesteps': Timesteps, 'PHASE_LABELS':
↪PHASE_LABELS,
            'FOCAL_PLANES': FOCAL_PLANES, 'NUM_PLANES': NUM_PLANES,
            'DDIM_STEPS': DDIM_STEPS, 'DDIM_ETA': DDIM_ETA, 'CONTEXT_K':
↪CONTEXT_K,
        }}, ckpt_path)
        print('Saved best to', ckpt_path)
    else:
        epochs_no_improve += 1

```



```

        print(f"No improvement for {epochs_no_improve} epoch(s)")

    # Generate samples after every N epochs
    print(f"\n--- Generating samples for epoch {epoch} ---")
    diffusion.eval()
    with torch.no_grad():
        for plane in tqdm_bar(SAMPLE_PLANES, desc=f"Epoch {epoch:03d} sample_
↳outputs"):
            sample_dir = OUT_SAMPLES_DIR / f'epoch_{epoch:03d}' / plane
            sample_dir.mkdir(parents=True, exist_ok=True)
            samples = val_dataset.sample_pairs_for_plane(plane,
↳SAMPLES_PER_PLANE)
            if samples:
                _save_sample_grid(samples, sample_dir, title_suffix='(mixed)')
                # Transitional-only samples
                trans_dir = sample_dir / 'transitions'
                trans_samples = val_dataset.
↳sample_transition_pairs_for_plane(plane, SAMPLES_PER_PLANE)
                if trans_samples:
                    _save_sample_grid(trans_samples, trans_dir,
↳title_suffix='(transition)')
                    # Rollout grids (full descriptive rollout of k steps)
                    rollout_dir = sample_dir / 'rollouts'
                    rollout_seed = val_dataset.sample_pairs_for_plane(plane, 1)
                    if rollout_seed:
                        _save_rollout_grid(rollout_seed[0], rollout_dir,
↳steps=CONTEXT_K)
                    print(f"Samples saved to {OUT_SAMPLES_DIR / f'epoch_{epoch:03d}'}\n")

    # Early stopping check
    if EARLY_STOPPING_PATIENCE and EARLY_STOPPING_PATIENCE > 0 and
↳epochs_no_improve >= EARLY_STOPPING_PATIENCE:
        print(f"Stopping early after {epochs_no_improve} epochs without_
↳improvement (patience={EARLY_STOPPING_PATIENCE}).")
        break

print(f"\nTraining complete! Best checkpoint saved to {ckpt_path}")

```

Epoch 001 Train: 100%| 14062/14062 [22:46<00:00, 10.29it/s]

Epoch 001 Val: 100%| 1889/1889 [02:05<00:00, 15.07it/s]

[Epoch 001] train 0.0820 | val 0.0643 | inv 0.0543 | SSIM 0.4915 | SSIM_tr
0.4744 | SSIM_non 0.4787 | LPIPS 0.1576 | LPIPS_tr 0.1539 | LPIPS_non 0.1488

Saved best to /mnt/c/Users/ioann/OneDrive/Documents/Projects/thesis/diffusion/em
bryo_diffusion_models/nextframe_ctxk7_128px.pt

--- Generating samples for epoch 1 ---

Epoch 001 sample outputs: 100%| 3/3 [00:16<00:00, 5.36s/it]
 Samples saved to /mnt/c/Users/ioann/OneDrive/Documents/Projects/thesis/diffusion
 /diffusion_samples/epoch_001

Epoch 002 Train: 100%| 14062/14062 [22:14<00:00, 10.54it/s]
 Epoch 002 Val: 100%| 1889/1889 [02:05<00:00, 15.03it/s]
 [Epoch 002] train 0.0609 | val 0.0579 | inv 0.0565 | SSIM 0.4600 | SSIM_tr
 0.4640 | SSIM_non 0.4617 | LPIPS 0.1698 | LPIPS_tr 0.1739 | LPIPS_non 0.1720
 Saved best to /mnt/c/Users/ioann/OneDrive/Documents/Projects/thesis/diffusion/em
 bryo_diffusion_models/nextframe_ctxk7_128px.pt

--- Generating samples for epoch 2 ---

Epoch 002 sample outputs: 100%| 3/3 [00:15<00:00, 5.09s/it]
 Samples saved to /mnt/c/Users/ioann/OneDrive/Documents/Projects/thesis/diffusion
 /diffusion_samples/epoch_002

Epoch 003 Train: 38%| 5280/14062 [08:30<09:30, 15.39it/s]

1.13 13. Training curves

Plot loss and validation metrics collected across epochs.

```
[ ]: # Plot training curves
if metric_history:
    epochs = np.arange(1, len(metric_history) + 1)
    train_loss = np.array([m['train_loss'] for m in metric_history],
    dtype=float)
    val_loss = np.array([m['val_loss'] for m in metric_history], dtype=float)
    val_ssim = np.array([m['val_ssim'] if m['val_ssim'] is not None else np.nan
    for m in metric_history], dtype=float)
    val_ssim_tr = np.array([m.get('val_ssim_tr') if m.get('val_ssim_tr') is not
    None else np.nan for m in metric_history], dtype=float)
    val_ssim_non = np.array([m.get('val_ssim_non') if m.get('val_ssim_non') is
    not None else np.nan for m in metric_history], dtype=float)
    val_lpips = np.array([m['val_lpips'] if m['val_lpips'] is not None else np.
    nan for m in metric_history], dtype=float)
    val_lpips_tr = np.array([m.get('val_lpips_tr') if m.get('val_lpips_tr') is
    not None else np.nan for m in metric_history], dtype=float)
    val_lpips_non = np.array([m.get('val_lpips_non') if m.get('val_lpips_non')
    is not None else np.nan for m in metric_history], dtype=float)

    fig, axes = plt.subplots(1, 3, figsize=(16, 4))
    axes[0].plot(epochs, train_loss, label='Train')
    axes[0].plot(epochs, val_loss, label='Val')
```

```


axes[0].set_title('Loss')
axes[0].set_xlabel('Epoch')
axes[0].set_ylabel('MSE')
axes[0].legend()
axes[0].grid(True, alpha=0.3)

axes[1].plot(epochs, val_ssim, label='Val SSIM')
axes[1].plot(epochs, val_ssim_tr, label='Val SSIM (transition)',
↳linestyle='--')
axes[1].plot(epochs, val_ssim_non, label='Val SSIM (non-transition)',
↳linestyle=':')
axes[1].set_title('SSIM')
axes[1].set_xlabel('Epoch')
axes[1].set_ylabel('SSIM')
axes[1].grid(True, alpha=0.3)
axes[1].legend()

axes[2].plot(epochs, val_lpips, label='Val LPIPS')
axes[2].plot(epochs, val_lpips_tr, label='Val LPIPS (transition)',
↳linestyle='--')
axes[2].plot(epochs, val_lpips_non, label='Val LPIPS (non-transition)',
↳linestyle=':')
axes[2].set_title('LPIPS')
axes[2].set_xlabel('Epoch')
axes[2].set_ylabel('LPIPS')
axes[2].grid(True, alpha=0.3)
axes[2].legend()

plt.tight_layout()
plt.show()
else:
    print('No metrics logged yet.')

```



1.14 14. Inference helpers

These functions generate ad-hoc samples or roll out multiple steps from a given context window.

```

[ ]: def _stack_context_paths(ctx_paths: list[str]) -> torch.Tensor:
    """Load k context images into a single [1,k,H,W] tensor."""
    assert len(ctx_paths) == CONTEXT_K, f"Expected {CONTEXT_K} context paths"
    ctx_tensors = [load_image(p).squeeze(0) for p in ctx_paths] # [k, H, W]
    return torch.stack(ctx_tensors, dim=0).unsqueeze(0) # [1, k, H, W]

def _infer_frame_index(path: str) -> int | None:
    """Infer frame index from filename like 00037.jpeg/jpg."""
    base = os.path.basename(path)
    m = re.search(r"(\d+)\.jpe?g$", base)
    return int(m.group(1)) if m else None

def _infer_time_scalar(context_paths: list[str], max_frame_index: int | None = None) -> float:
    # Use last context frame index as t_end; normalize by max_frame_index if
    # provided
    t_end = _infer_frame_index(context_paths[-1])
    if t_end is None:
        return 0.0
    if max_frame_index is None or max_frame_index <= 0:
        max_frame_index = t_end
    return float(t_end) / float(max(1, max_frame_index))

def visualize_samples(num=4, plane='F0', save_dir: str | None = None):
    """Save the same sample grids as training: mixed, transition, and rollout
    grids."""
    if save_dir is None:
        save_dir = OUT_SAMPLES_DIR / 'ad_hoc_samples' / plane
    save_dir = Path(save_dir)
    save_dir.mkdir(parents=True, exist_ok=True)

    # Mixed samples
    samples = val_dataset.sample_pairs_for_plane(plane, num)
    if samples:
        _save_sample_grid(samples, save_dir, title_suffix='(mixed)')

    # Transitional-only samples
    trans_dir = save_dir / 'transitions'
    trans_samples = val_dataset.sample_transition_pairs_for_plane(plane, num)
    if trans_samples:
        _save_sample_grid(trans_samples, trans_dir, title_suffix='(transition)')

    # Rollout grid (single seed)
    rollout_dir = save_dir / 'rollouts'
    rollout_seed = val_dataset.sample_pairs_for_plane(plane, 1)
    if rollout_seed:
        _save_rollout_grid(rollout_seed[0], rollout_dir, steps=CONTEXT_K)

```

```

    print(f"Ad hoc samples saved to {save_dir}")

def predict_next_frame(context_paths: list[str], out_path: str, max_frame_index:
    ↪ int | None = None):
    """Generate the next frame given k context frames."""
    assert len(context_paths) == CONTEXT_K, f"Need exactly {CONTEXT_K} context_
    ↪frames"
    x_ctx = _stack_context_paths(context_paths).to(device)
    frame_idx = None
    if USE_TIME_MAP:
        t_end = _infer_frame_index(context_paths[-1]) or 0
        frame_idx = min(int(t_end + 1), int(MAX_FRAME_INDEX))
        frame_idx = torch.tensor([frame_idx], device=device, dtype=torch.long)
    x_gen = sample_from_model(
        (1, 1, IMG_SIZE, IMG_SIZE),
        x_ctx=x_ctx,
        frame_idx=frame_idx,
    )
    out_path = str(out_path)
    save_tensor_image(x_gen[0], out_path)
    return out_path

def rollout_next_frames(
    initial_context_paths: list[str],
    steps: int,
    out_dir: str,
    max_frame_index: int | None = None
):
    """Multi-step rollout with a moving context window."""
    assert len(initial_context_paths) == CONTEXT_K, f"Need exactly {CONTEXT_K}
    ↪context frames"
    os.makedirs(out_dir, exist_ok=True)

    # initialize context tensors (each [1, H, W]) in [-1,1]
    ctx_tensors = [load_image(p) for p in initial_context_paths]
    t_end = _infer_frame_index(initial_context_paths[-1]) or 0

    generated_paths = []
    frame_idx = min(int(t_end + 1), int(MAX_FRAME_INDEX))
    for step in range(1, steps + 1):
        x_ctx = torch.stack([t.squeeze(0) for t in ctx_tensors], dim=0).
        ↪unsqueeze(0).to(device)
        if USE_TIME_MAP:
            frame_idx = min(frame_idx, int(MAX_FRAME_INDEX))
            frame_idx_tensor = torch.tensor([frame_idx], device=device,
            ↪dtype=torch.long)

```

```

else:
    frame_idx_tensor = None
    x_gen = sample_from_model(
        (1, 1, IMG_SIZE, IMG_SIZE),
        x_ctx=x_ctx,
        frame_idx=frame_idx_tensor,
    )
    save_path = os.path.join(out_dir, f'gen_t{t_end+1:05d}.png')
    save_tensor_image(x_gen[0], save_path)
    generated_paths.append(save_path)

    # update context window: drop oldest, append generated
    ctx_tensors = ctx_tensors[1:] + [x_gen[0].detach().cpu()]
    t_end += 1
    frame_idx += 1

return generated_paths

visualize_samples()

```

Ad hoc samples saved to /mnt/c/Users/ioann/OneDrive/Documents/Projects/thesis/diffusion/diffusion_samples/ad_hoc_samples/F0